

ML IN 2022

- * supervised Learning 99% by neural networks
 - self-supervised learning
 - Transfer learning
 - etc.
- * Unsupervised Learning
- * Reinforcement Learning 99% using neural networks

NEURAL NETWORKS

Artificial model of neuron (Perceptron) 1940's

↓
unnet promises 2 "AI winter"

A neural network

1960's

Feedforward Neural Network

Multi-layer Perceptron

Vanilla neural network

↓
unnet promises 2 better methods (SVM) 2 2nd AI winter

Convolutional Neural Network

2000's

* Big bang: AlexNet 2012

Tools: Caffe → TensorFlow
→ PyTorch

2012 - 2022 New Architectures

Backbones (AlexNet, VGGNet, ResNet, MobileNet, ...)

Encoder-Decoder

* U-Net

* DeepFake

GAN

NeRF

Transformer ⇒ Gen. purpose differentiable "computer"

SUPERVISED ML

Essentially 3 main problems that are diff evaluated:

1. Classification

* Assign one of K classes

* Perf. meas. classification percent/rate 0...1.0

2. Regression

* Estimate output given inputs

* Perf. by MSE / MAE

3. Detection

* Tell if (and when) event X present or not

* Perf. by precision-recall (ROC curve)

\Rightarrow two types of errors FP & FN

LINEAR MODEL

For regression:

x : input

y : output


a, b : tr. parameters
aka weights

0: points

1 point

2 points \Rightarrow high school

N points \Rightarrow Least Squares
Fit



$$\begin{cases} y_1 = ax_1 + b \Rightarrow b = y_1 - ax_1 \\ y_2 = ax_2 + b \Rightarrow y_2 = ax_2 + y_1 - ax_1 \end{cases}$$

$$\Rightarrow \boxed{a = \frac{y_2 - y_1}{x_2 - x_1}}$$

For classification:

x_1, \dots, x_3 from class 1


x_4, \dots, x_6 from class 2

y ?

Let's assign

$y = -1$ for C_1

$y = +1$ for C_2



\Rightarrow Fit line

$$\text{class} = \begin{cases} 1 & \text{if } ax+b < 0 \\ 2 & \text{if } ax+b > 0 \end{cases}$$

Problem: Line fitting error depends on your distance to the discriminator point $y=0$ while intuitively $\text{err}=0$ for correct and $\text{err}=1$ for wrong classification should be used.

○

Fix:

$y = 0$ for C_1


$y = 1$ for C_2

Now ground truth and prediction error

$$\|y_{\text{GT}} - \hat{y}\| = \{0, 1\}$$

or

$$(y_{\text{GT}} - \hat{y})^2 = \begin{cases} 0 & \text{if } y_{\text{GT}} = \hat{y} \\ 1 & \text{if } y_{\text{GT}} \neq \hat{y} \end{cases}$$



step function is approximated by sigmoid

$$\text{logsig}(x) = \frac{1}{1 + e^{-x}}$$

Linear classifier:

$$y = \text{logsig}(ax + b)$$

Optimization?

- (1) No closed form solution, but we can minimize error over N training samples

$$\begin{aligned} \underset{\uparrow}{\mathcal{L}_{\text{MSE}}} &= \frac{1}{N} \sum_{n=1}^N (y_{\text{GT}}^{(n)} - \hat{y}_n)^2 = \frac{1}{N} \sum_{n=1}^N (y_{\text{GT}}^{(n)} - ax_n - b)^2 \quad \text{reg.} \\ &= \frac{1}{N} \sum_{n=1}^N (y_{\text{GT}}^{(n)} - \text{logsig}(ax_n + b))^2 \quad \text{class.} \end{aligned}$$

"loss" in ML terminology

We can go toward (local) minimum by taking a small step μ (learning rate)

- (2) toward negative gradient.

$$a^{t+1} = a^t - \mu \frac{\partial \mathcal{L}_{\text{MSE}}}{\partial a}$$

$$b^{t+1} = b^t - \mu \frac{\partial \mathcal{L}_{\text{MSE}}}{\partial b}$$

\Rightarrow Gradient Descent (GD) algorithm

Start from some random point (a^0, b^0)

Stuck to local minima?

$$\frac{\partial L_{MSE}}{\partial a} = \frac{\partial}{\partial a} \frac{1}{N} \sum_{n=1}^N (y_{GT}^{(n)} - \text{logsig}(ax_n + b))^2$$

$$\left[\frac{\partial}{\partial x} f^2(x) = 2f(x) \cdot f'(x) \quad \text{chain rule} \right]$$

$$= \frac{2}{N} \sum_{n=1}^N (y_{GT}^{(n)} - \text{logsig}(ax_n + b)) \cdot \frac{\partial}{\partial a} (y_{GT}^{(n)} - \text{logsig}(ax_n + b))$$

$$\left[\frac{\partial}{\partial x} \text{logsig}(x) = \text{logsig}(x) (1 - \text{logsig}(x)) \right]$$

$$= \frac{2}{N} \sum_{n=1}^N (y_{GT}^{(n)} - \text{logsig}(ax_n + b)) \cdot -\text{logsig}(ax_n + b) \cdot (1 - \text{logsig}(ax_n + b))$$


$$\cdot \underbrace{\frac{\partial}{\partial a} (ax_n + b)}_{= x_n}$$

\Rightarrow it's annoyingly long, but CPU does that for us once it is programmed!

A SINGLE NEURON

Biological neuron : <NOTEBOOK>

Artificial models:



$$y' = w_1 \cdot x_1 + w_2 \cdot x_2 - b$$

Single neuron is a linear model that can be trained using gradient descent.

<NOTEBOOK>