

RECAP SO FAR

* model of a single neuron (perceptron)




* Gradient descent learning (GD)

* Multi-layer Perceptron (MLP)

* GD for MLP

⇒ can learn anything from sufficient data




For 100x100 image w/ 100 hidden layer-neurons + 3 output neurons

hidden weights

$$\underbrace{100 \times 100 \times (100 + 1)}_{\text{Flatten}} + (100 + 1) \times 3 = 1\text{M weights}$$

Decrease # of weights 1: Localized detectors




one location

$$\underbrace{25 \times 25 \times (25 + 1)}_{\text{one location}} \times 4 + (100 + 1) \times 3 = 65\text{k weights}$$

Prob.


lack of invariance



vs.




Extension to solution 1: shared weights



$$25 \times 25 \times (25+1) + (100+1) \times 3 = 16k \text{ weights}$$

By calculating extension in all locations we essentially have a convolutional neural network (CNN).

CNN Learning Tricks




every location
produces gradient and
only few are meaningful
 \Rightarrow a lot of noise

Solution 1: stride

We jump over every N th pixel (stride)


Stride 2



Solution 2: (Max) Pooling

Only the maximum response in certain neighborhood is retained

2x2 max pooling



After stride and pooling only 6.25% of the locations produce gradient update. In practice even less as only high response channels are selected \Rightarrow Filters specialize


Lower levels learn elementary features and higher levels more advanced feature (similar to human visual system)

[Video: Deep Dreams]

[notebook]

MNIST - Digits 0-9

Tested architectures



model =

`tf.keras.models.Sequential()`

* add layers one by one

`model.compile()`



* makes computation graph

`model.fit()`

* training (batch-size, learning-rate)

`model.predict()`

* use model



+ multiple conv + max pooling [homework]